

ECE284 Fall 21 W1S2

Low-power VLSI Implementation for Machine Learning

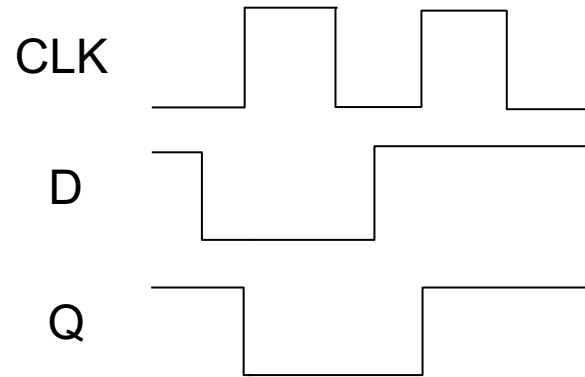
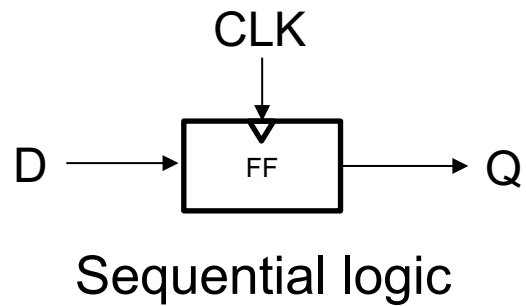
Prof. Mingu Kang

UCSD Computer Engineering

Announcement

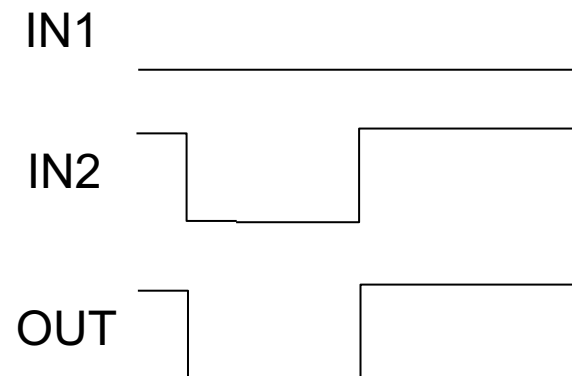
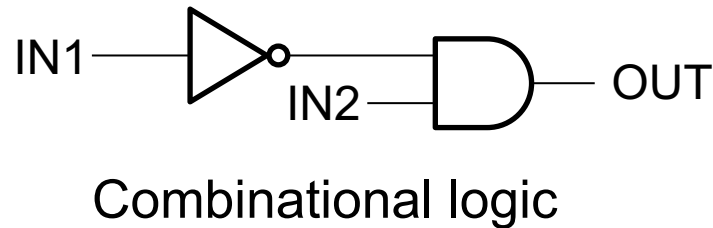
1. If you just enrolled, please visit “Canvas – Pages – In-class Video” tab.
2. If you just enrolled, please complete three modules.
3. Enrollment issue. Please contact me. ece284ucsd@gmail.com
4. Linuxcloud & Datahub check.
5. Office hour change:
 - Instructor: Friday 11 – 11:50 am
 - TA: Friday 3 – 3:50 pm

Verilog Review



```
reg Q;  
reg D;
```

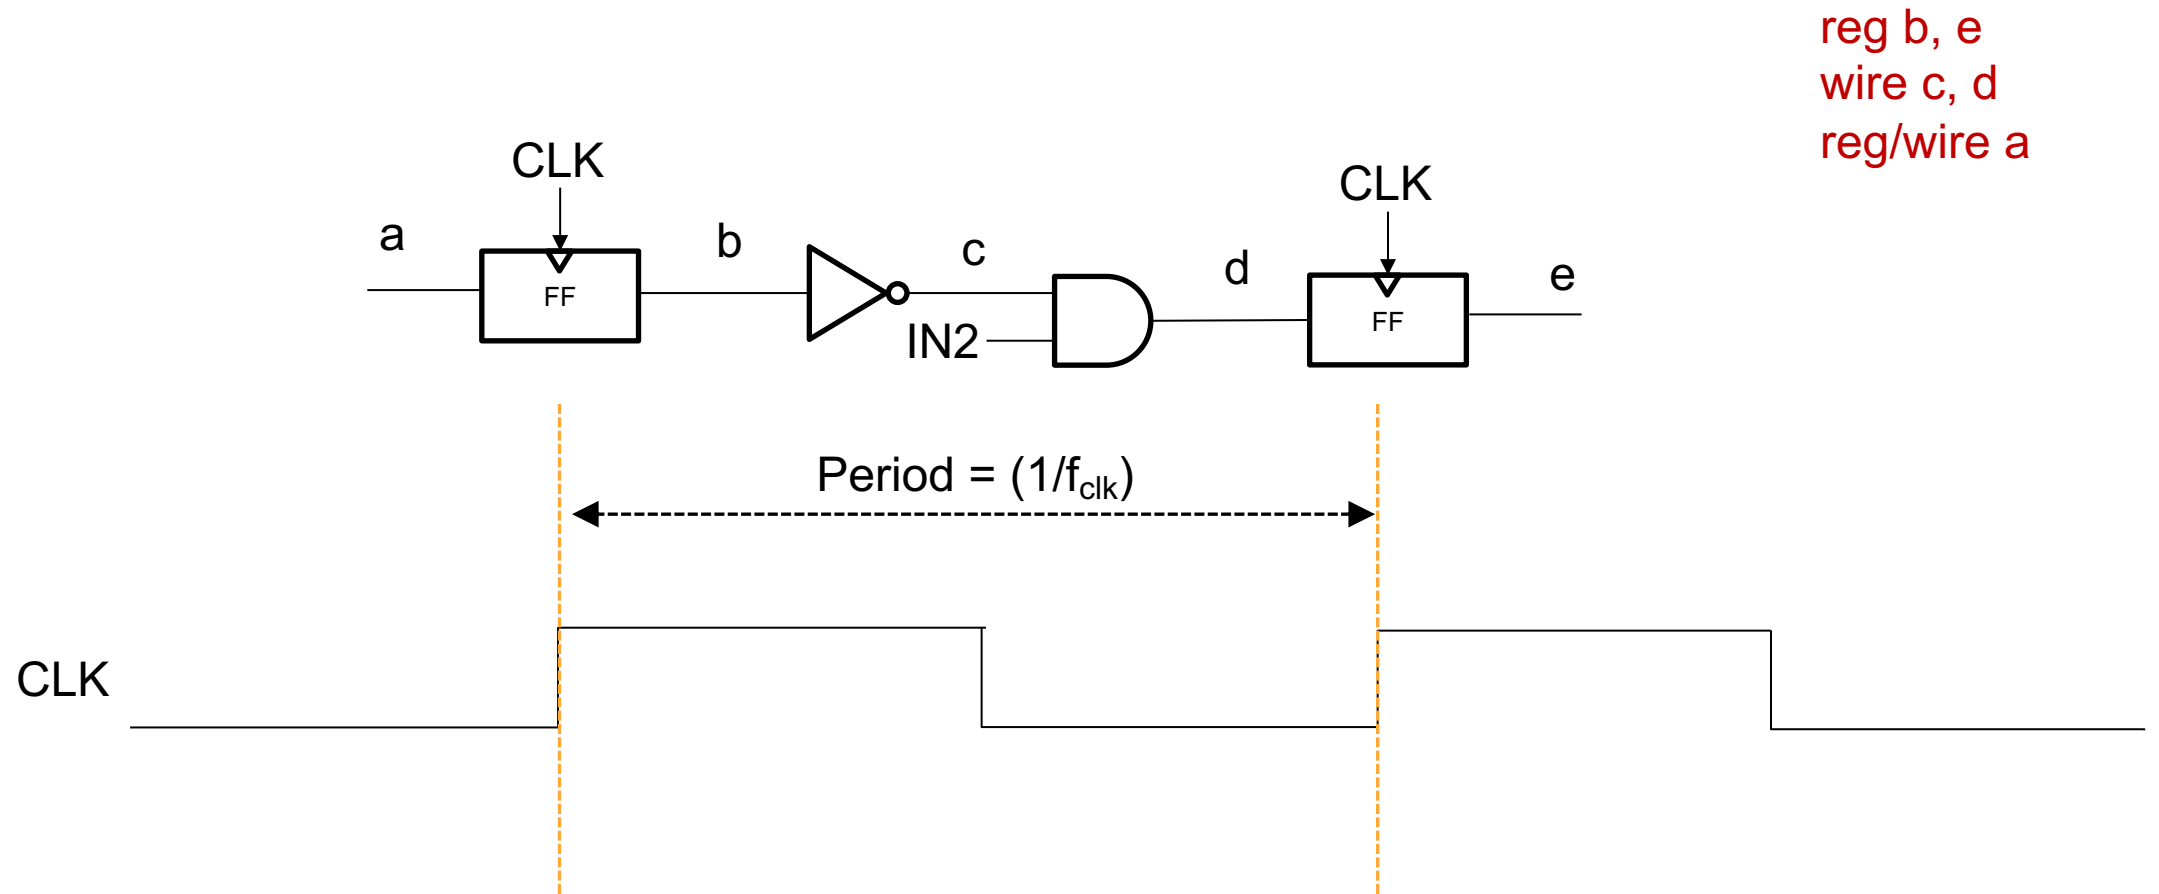
```
always @ (posedge clk) begin  
    Q <= D;  
end
```



```
wire IN1;  
wire IN2;  
wire OUT;
```

```
assign OUT = (!IN1) && (IN2);
```

Verilog Review



Verilog Review: Frequently used patterns

```
module risc (clk, add, sub, a, b, out);
```

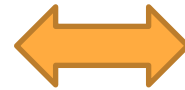
```
    input add, sub, clk;  
    input [3:0] a;  
    input [3:0] b;  
    output [7:0] out;
```

```
    reg out_q;
```

```
    assign out = out_q;
```

```
    always @ (posedge clk) begin  
        if (add)  
            out_q <= a+b;  
        else if (sub)  
            out_q <= a-b;  
    end
```

```
endmodule
```



```
module risc (clk, add, sub, a, b, out);
```

```
    input add, sub, clk;  
    input [3:0] a;  
    input [3:0] b;  
    output reg [7:0] out;
```

```
    always @ (posedge clk) begin  
        if (add)  
            out <= a+b;  
        else if (sub)  
            out <= a-b;  
    end
```

```
endmodule
```

Verilog Review: Frequently used patterns

```
module risc (clk, add, sub, a, b, out);
```

```
    input add, sub, clk;  
    input [3:0] a;  
    input [3:0] b;  
    output reg [7:0] out;
```

```
    always @ (posedge clk) begin  
        if (add)  
            out <= a+b;  
        else if (sub)  
            out <= a-b;  
    end
```

```
endmodule
```



```
module risc (clk, add, sub, a, b, out);
```

```
    input add, sub, clk;  
    input [3:0] a;  
    input [3:0] b;  
    output reg [7:0] out;
```

```
    assign result = (add)? a+b : ((sub)? a-b : result);
```

```
    always @ (posedge clk) begin  
        out <= result;  
    end
```

```
endmodule
```

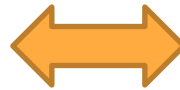
Verilog Review: Frequently used patterns

```
module risc (clk, inst, a, b, out);
```

```
    input inst;  
    input clk;  
    input [3:0] a;  
    input [3:0] b;  
    output reg [7:0] out;
```

```
    always @ (posedge clk) begin  
        case (inst)  
            1'b0: out <= a+b;  
            1'b1: out <= a-b;  
        endcase  
    end
```

```
endmodule
```



```
module risc (clk, inst, a, b, out);
```

```
    parameter bw = 4;
```

```
    input inst;  
    input clk;  
    input [bw-1:0] a;  
    input [bw-1:0] b;  
    output reg [2*bw-1:0] out;
```

```
    always @ (posedge clk) begin  
        case (inst)  
            1'b0: out <= a+b;  
            1'b1: out <= a-b;  
        endcase  
    end
```

```
endmodule
```

Verilog Review: Frequently used patterns

```
module risc (clk, a, b, out);
```

```
    parameter bw = 4;
```

```
    input clk;
```

```
    input [2*bw-1:0] a;
```

```
    input [2*bw-1:0] b;
```

```
    output reg [2*bw-1:0] out;
```

```
    wire [2*bw-1:0] out_temp;
```

```
    genvar i;
```

```
    for (i = 0; i < 2; i=i+1) begin
```

```
        assign out_temp[ bw*(i+1)-1 : bw*i ] = a[bw*(i+1)-1 : bw*i ] * b[bw*(i+1)-1 : bw*i];
```

```
    end
```

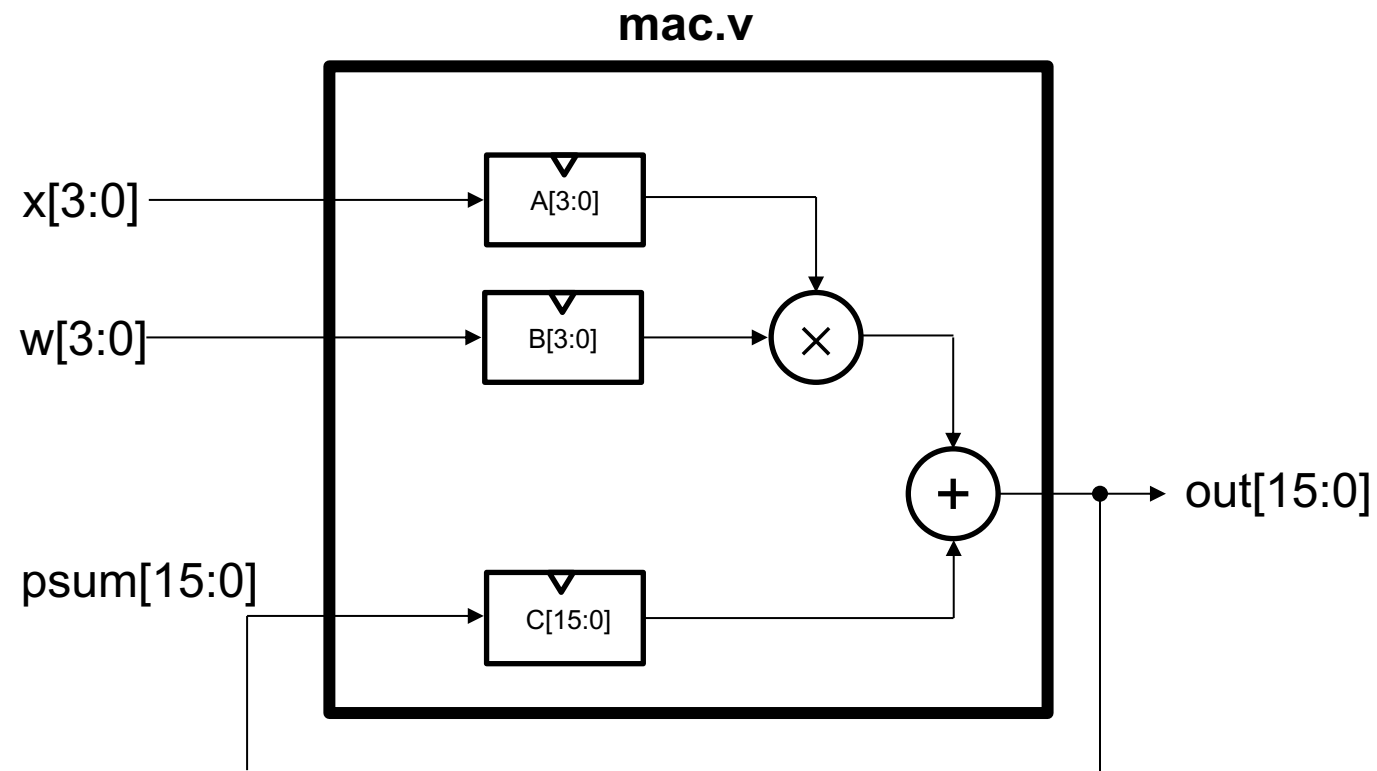
```
    always @ (posedge clk)
```

```
        out <= out_temp;
```

```
endmodule
```

<pre>out[7:0] = a[3:0]*b[3:0] out[15:8] = a[7:4]*b[7:4]</pre>
--

[Verilog] MAC Example



Alias

- **iveri** = 'iverilog -o compiled -c'
- **irun** = 'vvp compiled'
- **wave** = 'gtkwave'

Example commands

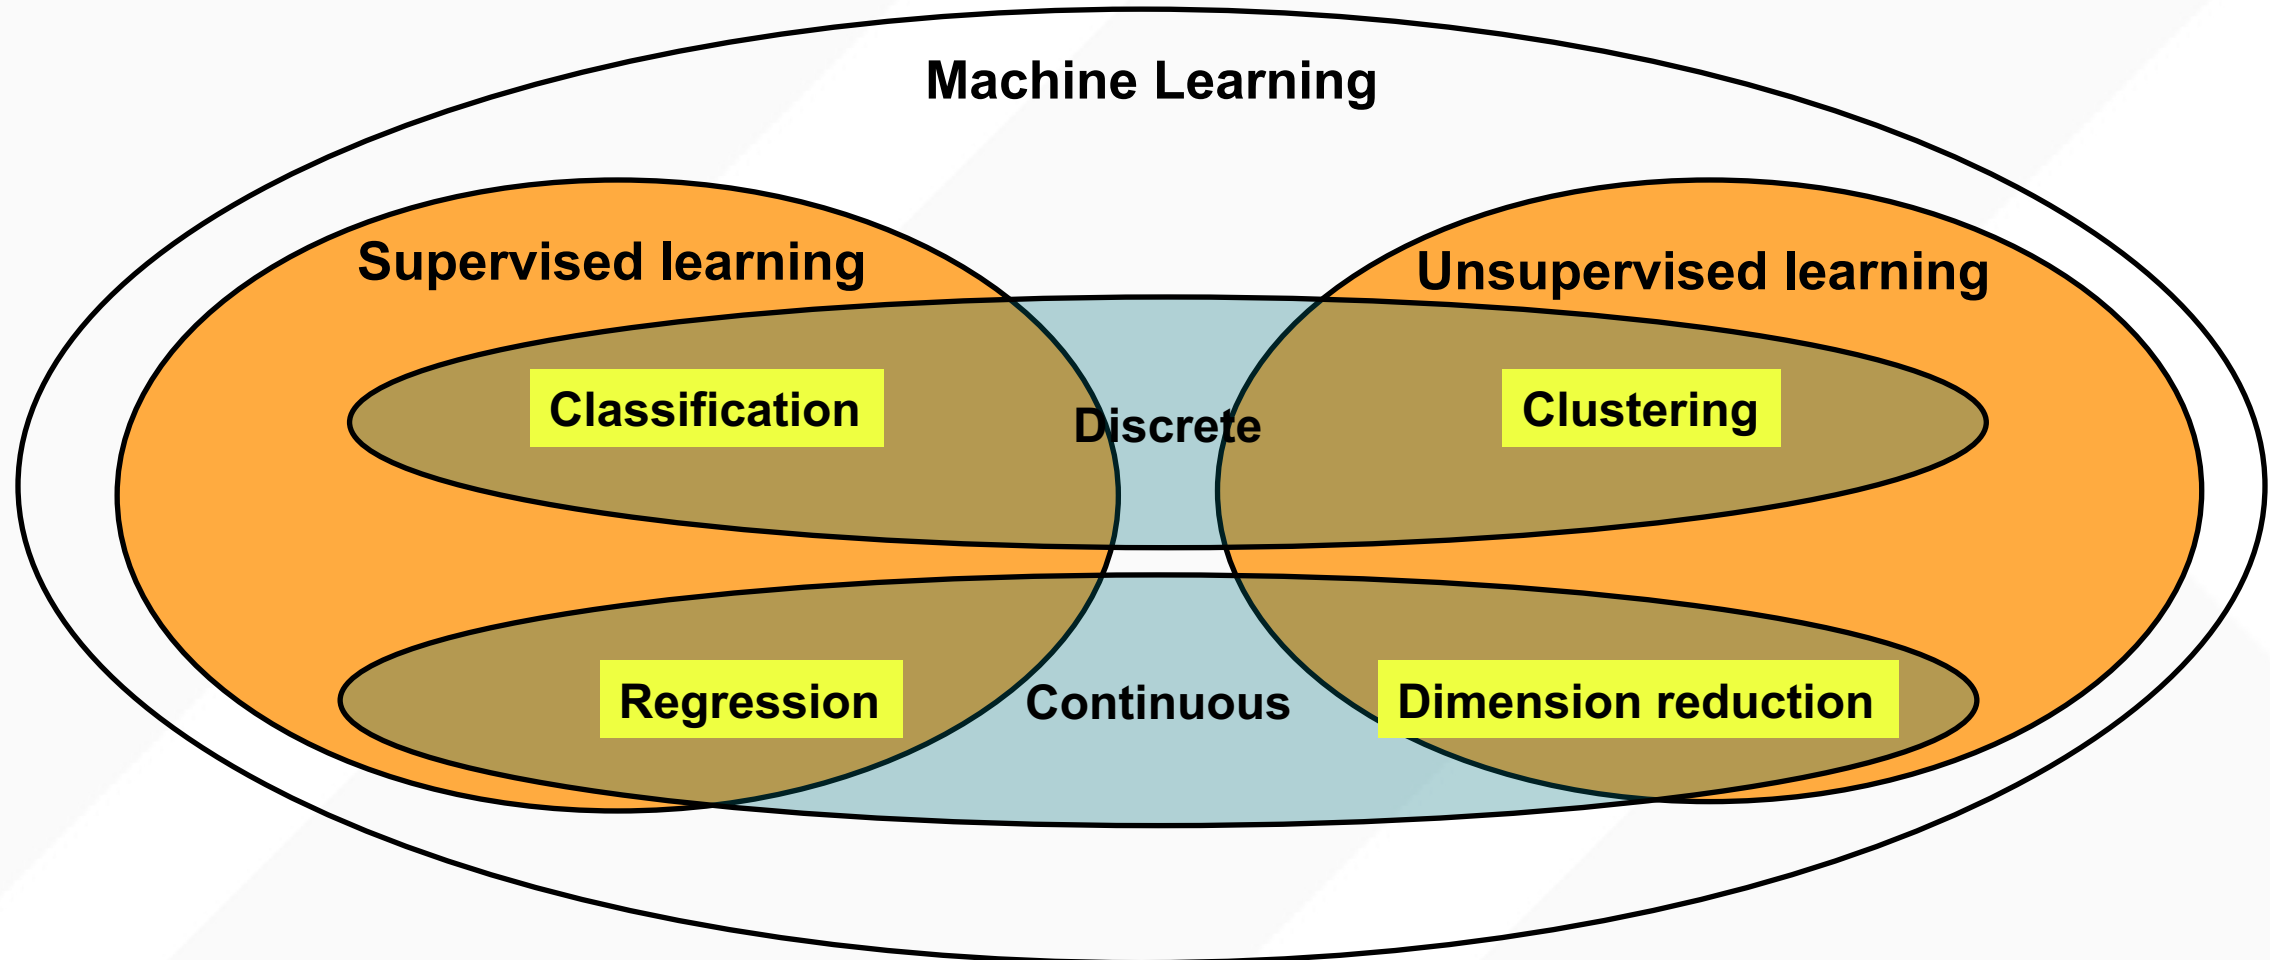
1. iveri filelist
(filelist includes all *.v files)
2. irun
3. wave mac_tb.vcd

- Weight / Activation: 4-bit, psum / output: 16-bit
- Weight, activations, psum are latched, but output is not latched
- Weight and activation data are in b_data.txt and a_data.txt files
- Output is fed back to psum input

[HW1] SFP Design

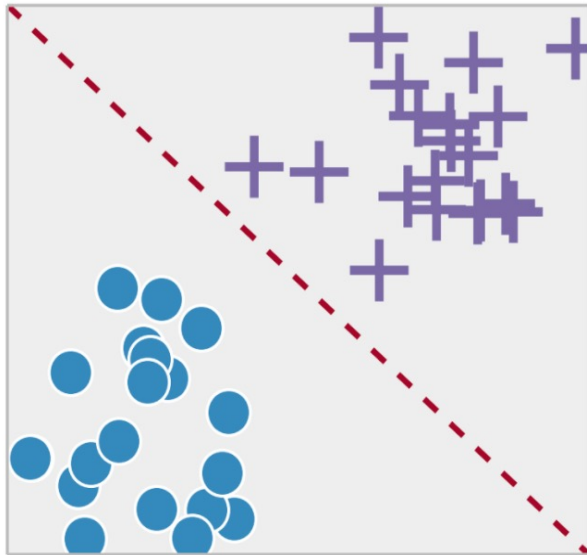
- In: 4-bit, out: 16-bit
- Control bits: acc, relu, reset
- If reset, internal latch "psum_q" becomes zero
- If $\text{acc} == 1$, psum_q will be updated with " $\text{psum_q} + \text{In}$ " in the next rising edge
- If $\text{relu} == 1$, psum_q is negative number, psum_q will be updated to be zero in the next rising edge
- out port is just connected to the psum_q
- Sample vcd file is attached in git.

Types of Machine Learning



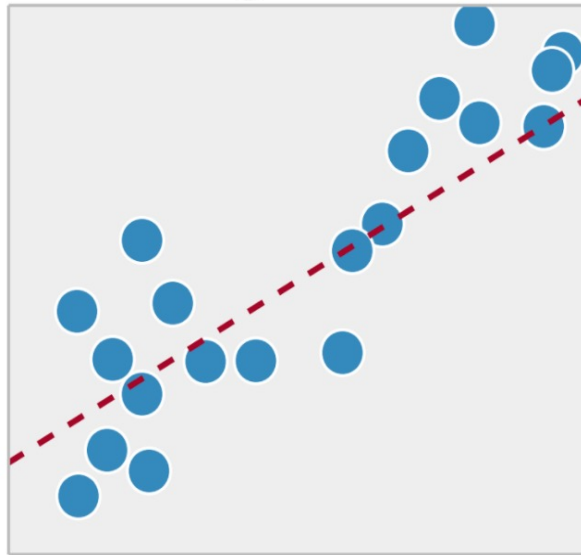
Types of Machine Learning – cont.

Classification



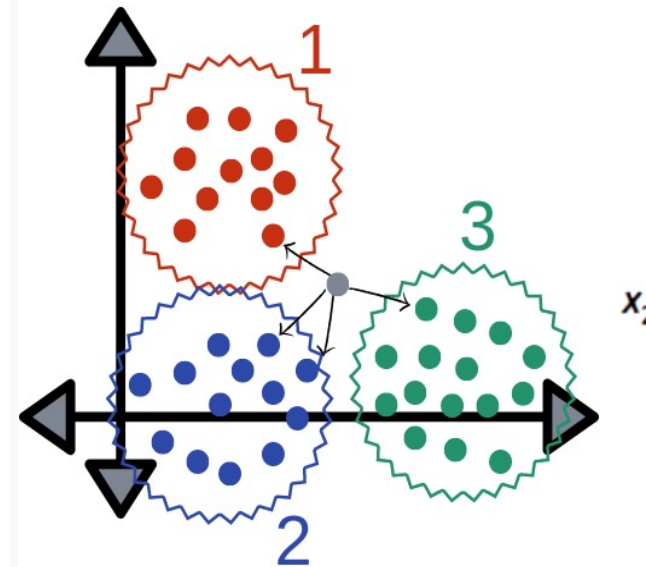
- Neural network
- Support vector machine
- ...

Regression



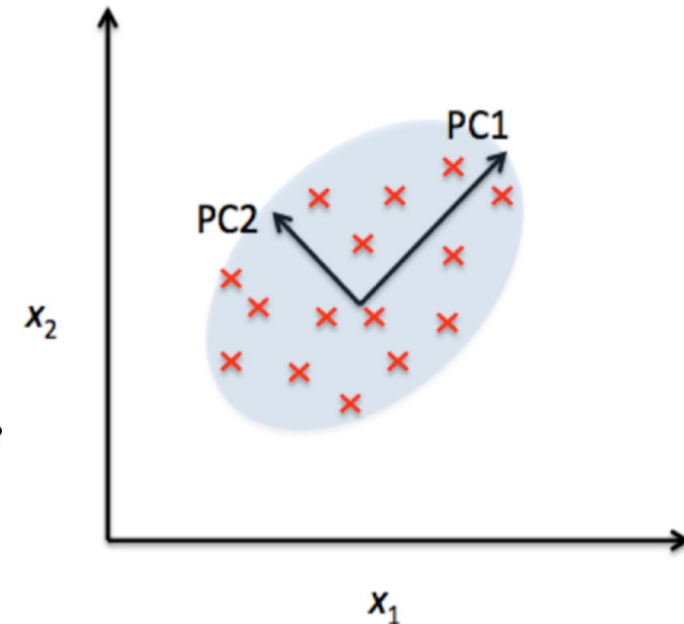
- Linear regression
- Polynomial Regression
- ...

Clustering



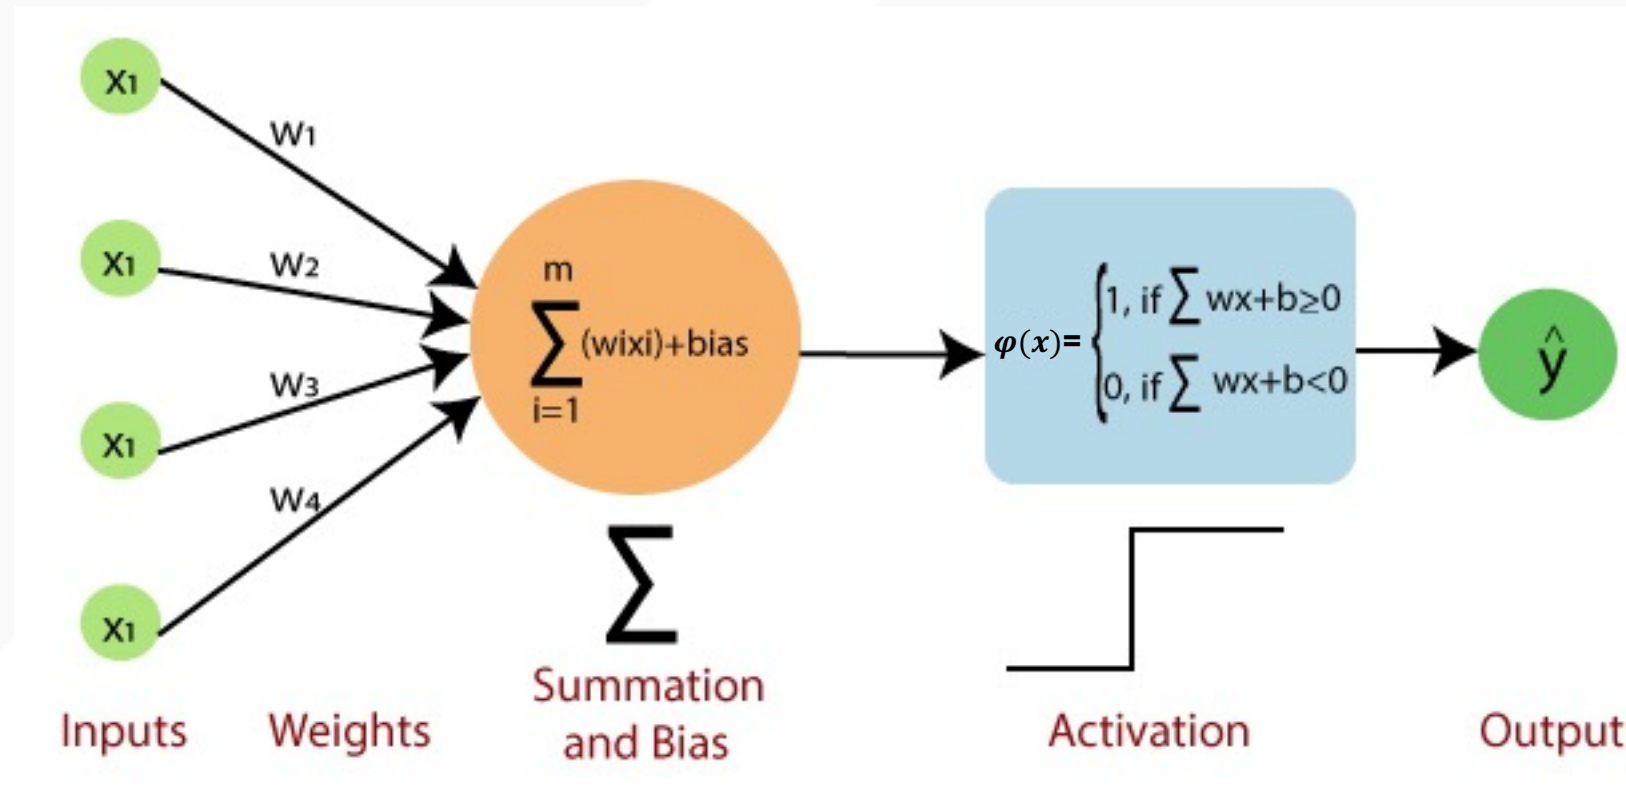
- K-nearest neighbor
- ...

Dimension reduction



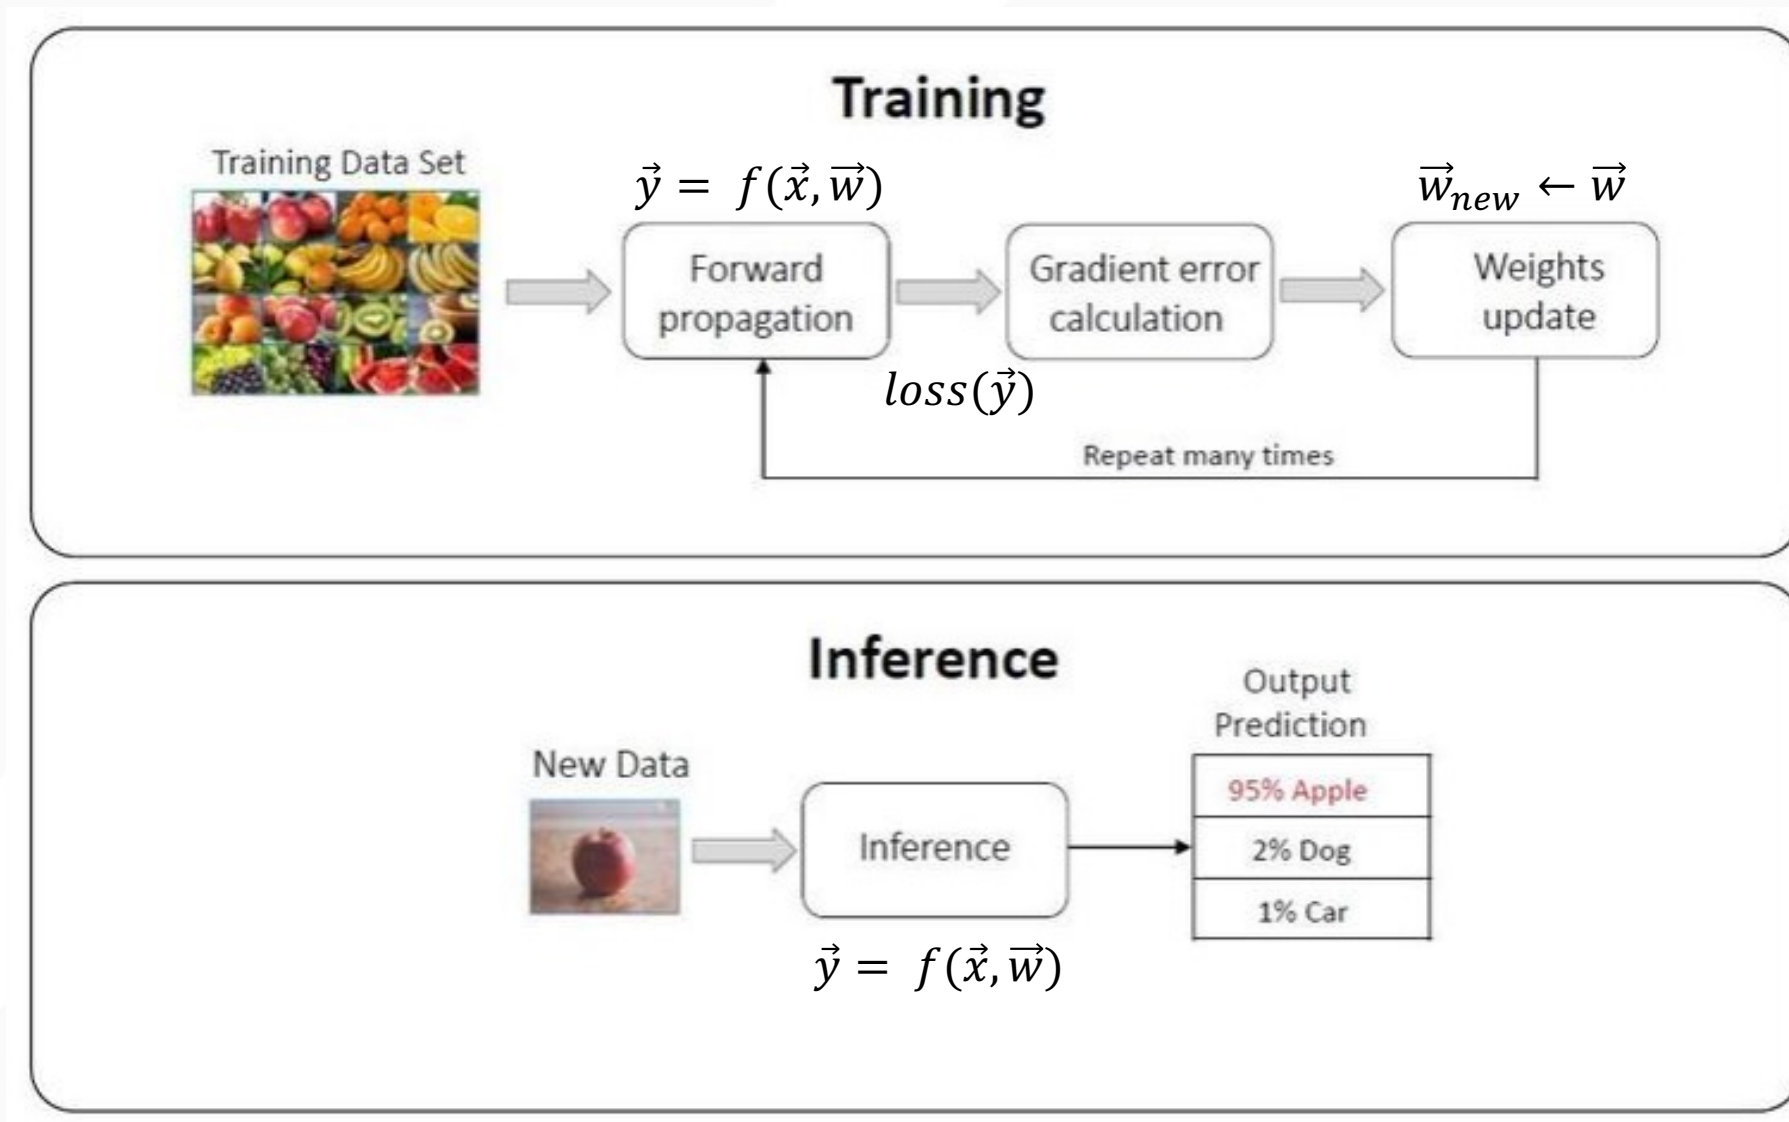
- Principle component analysis (PCA)
- Linear discriminant analysis (LDA)
- ...

Example of Machine Learning (Inference): Perceptron

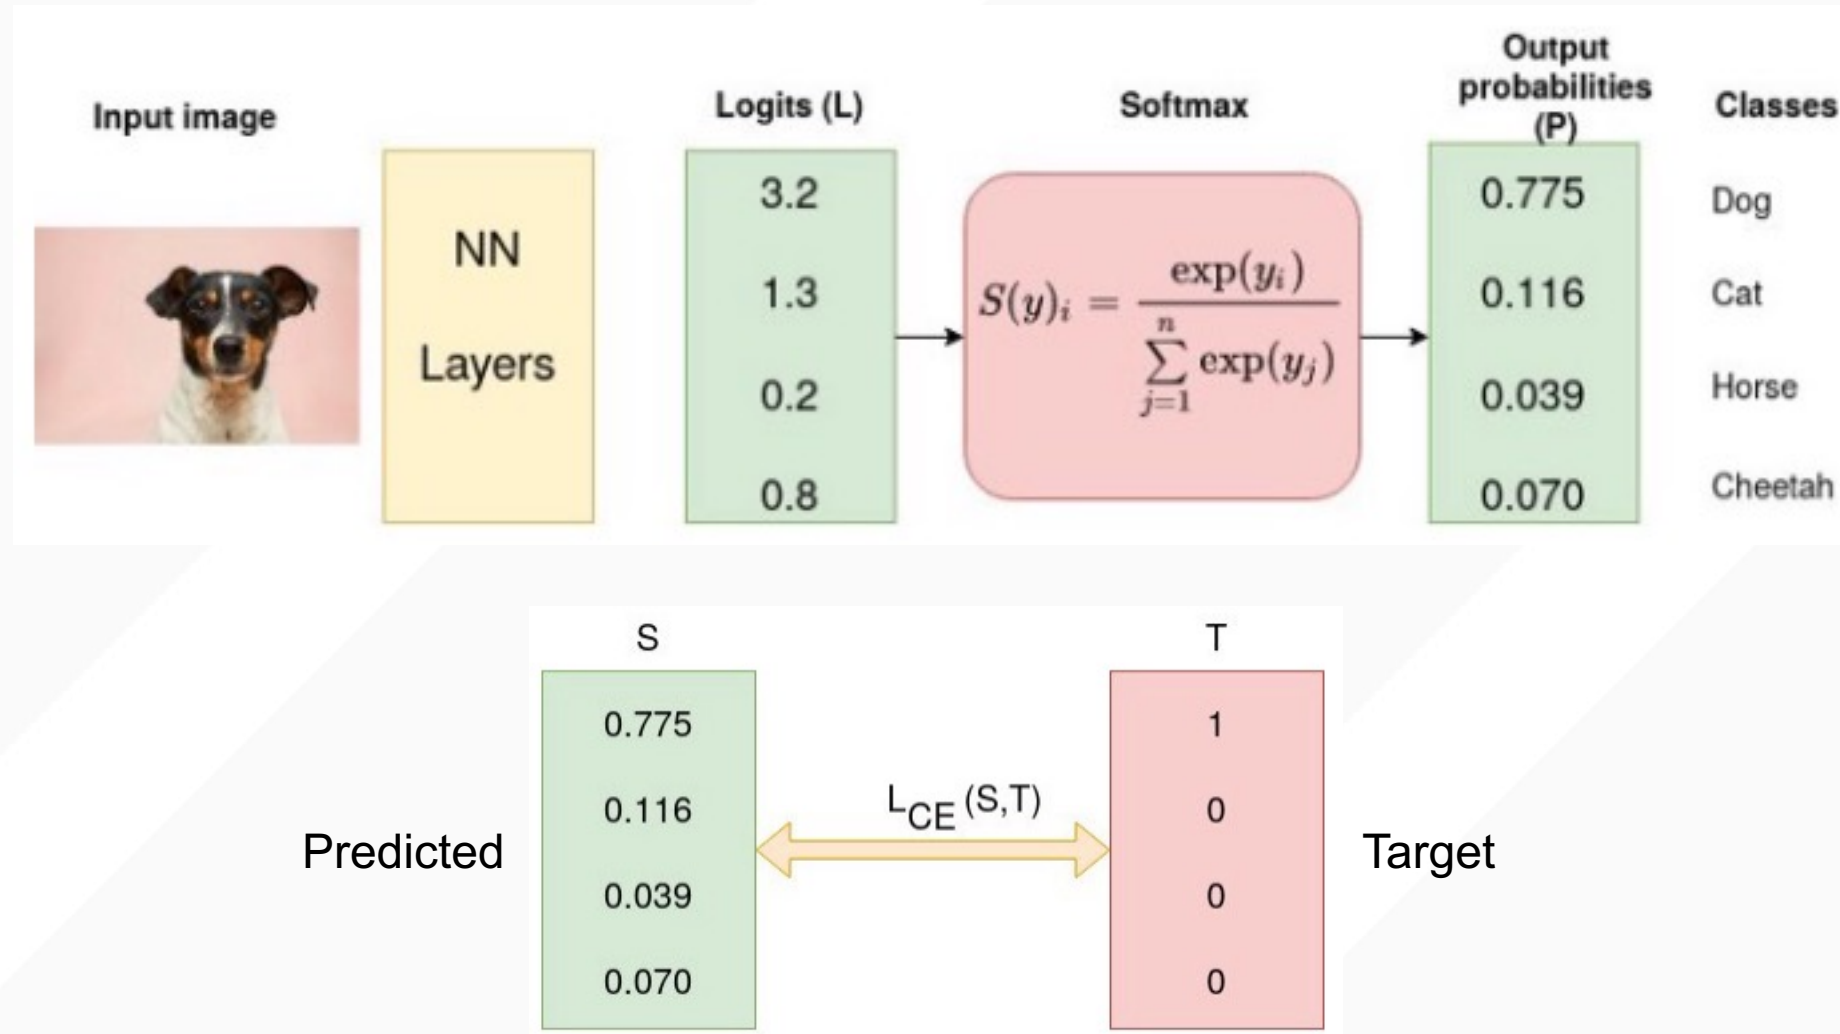


$$y = \varphi \left(\sum_{i=1}^n w_i x_i + b \right) = \varphi(w^T x + b)$$

Machine Learning – Training vs. Inference



Training



Loss (Cost) Function (Mean Squared Error)

Mean squared error loss

```
1 # calculate mean squared error
2 def mean_squared_error(actual, predicted):
3     sum_square_error = 0.0
4
5     for i in range(len(actual)):
6         sum_square_error += (actual[i] - predicted[i])**2.0
7     mean_square_error = sum_square_error / len(actual)
8     return mean_square_error
```

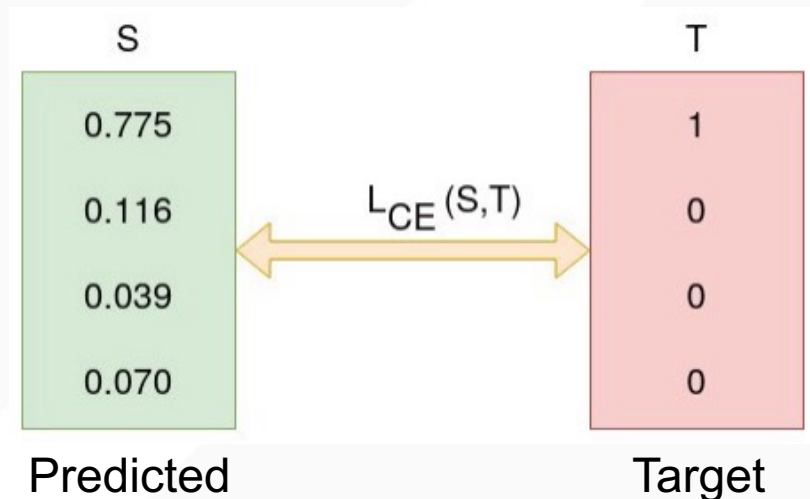

Loss (Cost) Function (Cross-Entropy Loss)

**Cross-Entropy loss
(or logarithmic loss,
or logistic loss)**

```

1 # calculate cross entropy
2 def categorical_cross_entropy(actual, predicted):
3     sum_score = 0.0
4
5     for i in range(len(actual)):           # number of class
6         sum_score += actual[i] * log(1e-15 + predicted[i])
7     mean_sum_score = 1.0 / len(actual) * sum_score
8     return -mean_sum_score
9

```



$$\begin{aligned}
 L_{CE} &= - \sum_{i=1} T_i \log(S_i) \\
 &= - [1 \log_2(0.775) + 0 \log_2(0.126) + 0 \log_2(0.039) + 0 \log_2(0.070)] \\
 &= - \log_2(0.775) \\
 &= 0.3677
 \end{aligned}$$

figures: <https://towardsdatascience.com/cross-entropy-loss-function-f38c4ec8643e>

<https://machinelearningmastery.com/loss-and-loss-functions-for-training-deep-learning-neural-networks/>

Gradient Descent (GD) Algorithm

```

1  # calculate cross entropy
2  def gradient_descent(gradient, start, learn_rate, n_iter):
3      vector = start
4      for _ in range(n_iter):
5          diff = -learn_rate * gradient(vector)
6          vector += diff
7      return vector

```

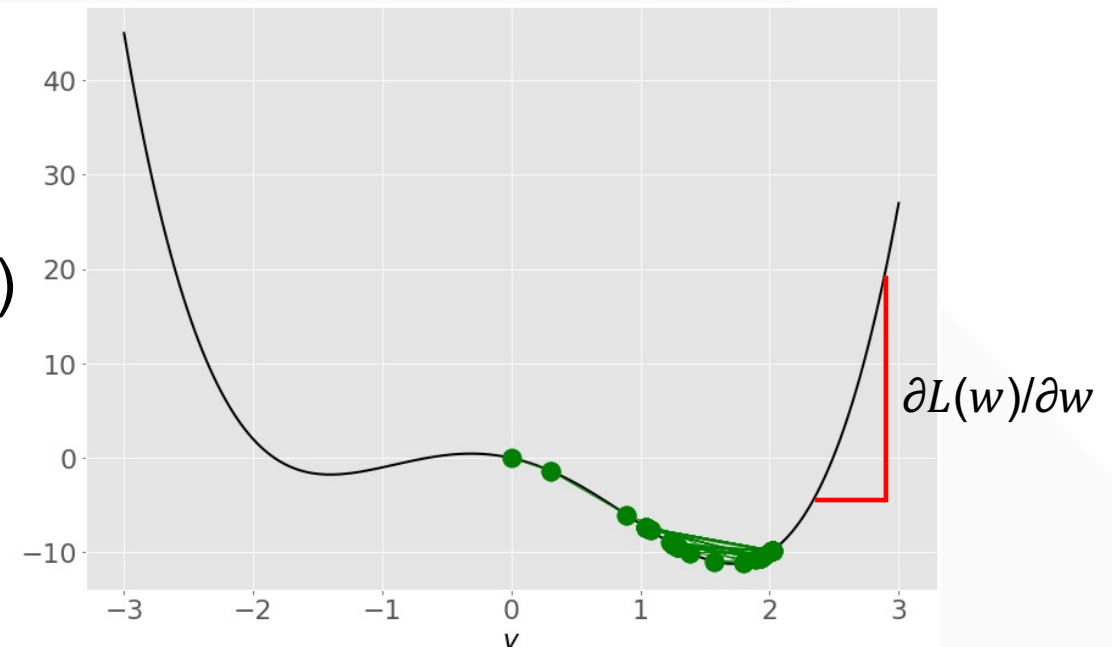
Finding optimal \mathbf{v} to minimize cost L

$$\mathbf{w} = (w_1, \dots, w_n)$$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla L(\mathbf{w})$$

- gradient $\nabla L(\mathbf{w}) = (\partial L / \partial w_1, \dots, \partial L / \partial w_n)$
- learning rate η

$L(w)$

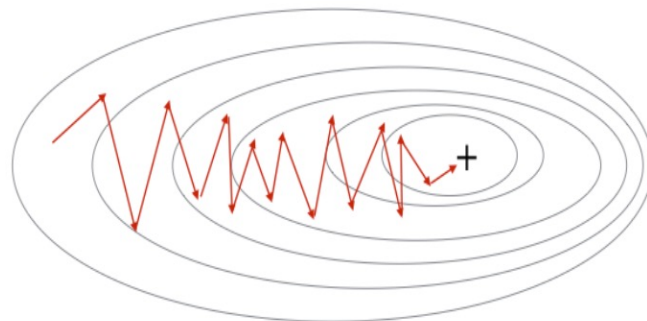


$L(w)$ vs. 1-dim w

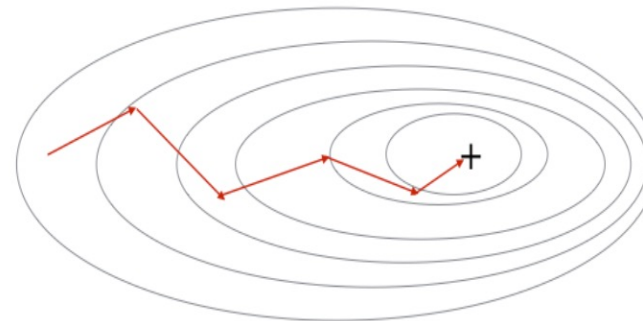
Variations of Gradient Decent Algorithms

- (Batch) Gradient Descent algorithm:
 - Every updated, the gradient computed with all the data points $\sum_{i \in all} \nabla_i$
 - Computation is slow, but update is smooth
- Stochastic Gradient Descent (SGD) algorithm:
 - Every updated, the gradient computed with only single data point
 - Computation is fast, but the training is noisy
- Mini-batch Gradient Descent (SGD) algorithm:
 - Every updated, the gradient computed with $\sum_{i \in subset} \nabla_i$

Stochastic Gradient Descent



Mini-Batch Gradient Descent



Calculation of Gradient with Mean Square Error Example

$$L(w, b) = MSE = \frac{1}{m} \sum_{j=1}^m (y_j - \hat{y}_j)^2 \quad \hat{y} = \phi(wx + b), \quad \phi(x) = x$$

, where m data points are used to train

$$\nabla L(w) = \frac{\partial L}{\partial w} = \frac{1}{m} \sum_{j=1}^m (y_j - \hat{y}_j) * (-2x_j)$$

$$\nabla L(b) = \frac{\partial L}{\partial b} = \frac{1}{m} \sum_{j=1}^m (y_j - \hat{y}_j) * (-2)$$

What if $\hat{y} = \phi(\sum_{i=1}^n w_i x_i + b)$, $\phi(x) = x$?

$$\nabla L(w_i) = \frac{\partial L}{\partial w_i} = \frac{1}{m} \sum_{j=1}^m (y_j - \hat{y}_j) * (-2x_{ji}), \text{ here } i = 1, 2, \dots, n$$

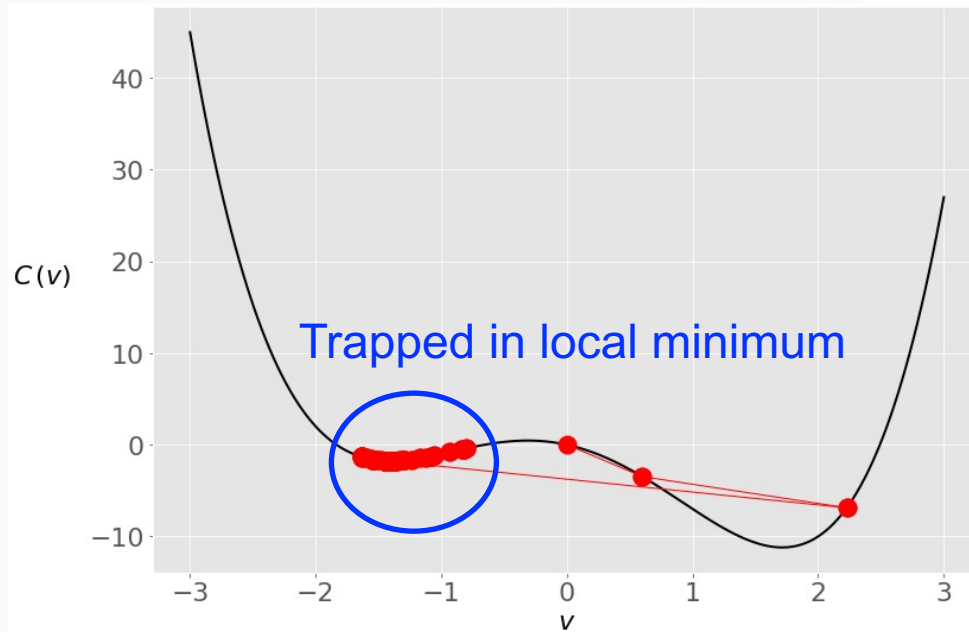
Type of Activations

Function Type	Equation	Derivative
Linear	$f(x) = ax + c$	$f'(x) = a$
Sigmoid	$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x) (1 - f(x))$
TanH	$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ReLU	$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parametric ReLU	$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
ELU	$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$

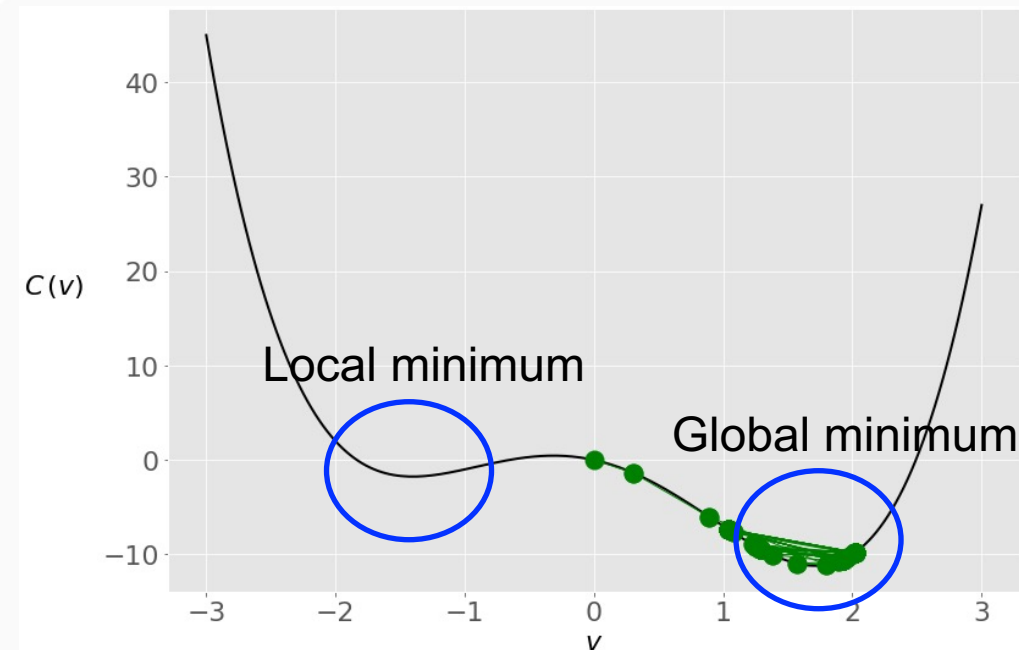
Gradient Descent (Example1)

https://pytorch.org/tutorials/beginner/pytorch_with_examples.html

Impact of Learning Rate (η)



Bad training



Desirable training

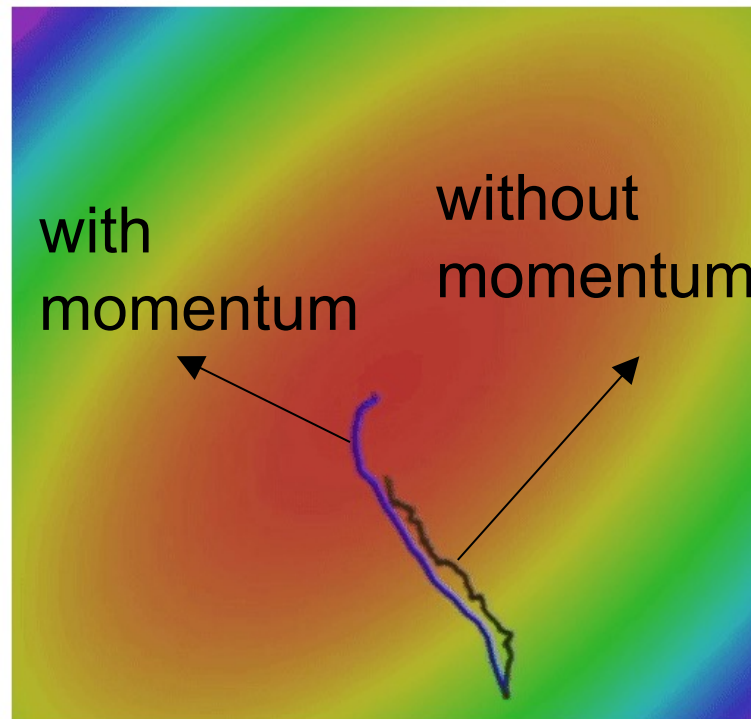
- Large learning rate: - could skip the global minimum
- might not converge, but oscillates
- Small learning rate: training is very slow

Learning Rate Decaying

```
1 # calculate cross entropy
2 def gradient_descent(gradient, start, learn_rate, n_iter):
3     vector = start
4     initial_learn_rate = large number
5     for _ in range(n_iter):
6         diff = - learn_rate * gradient(vector)
7         vector += diff
8         learn_rate = initial_learn_rate * (1 / (1 + decay * iteration))
    return vector
```

- Learning rate decaying:
 - Initially start with large learning rate for fast learning
 - Learning rate decreases as iteration goes on for better convergence

Momentum (v)

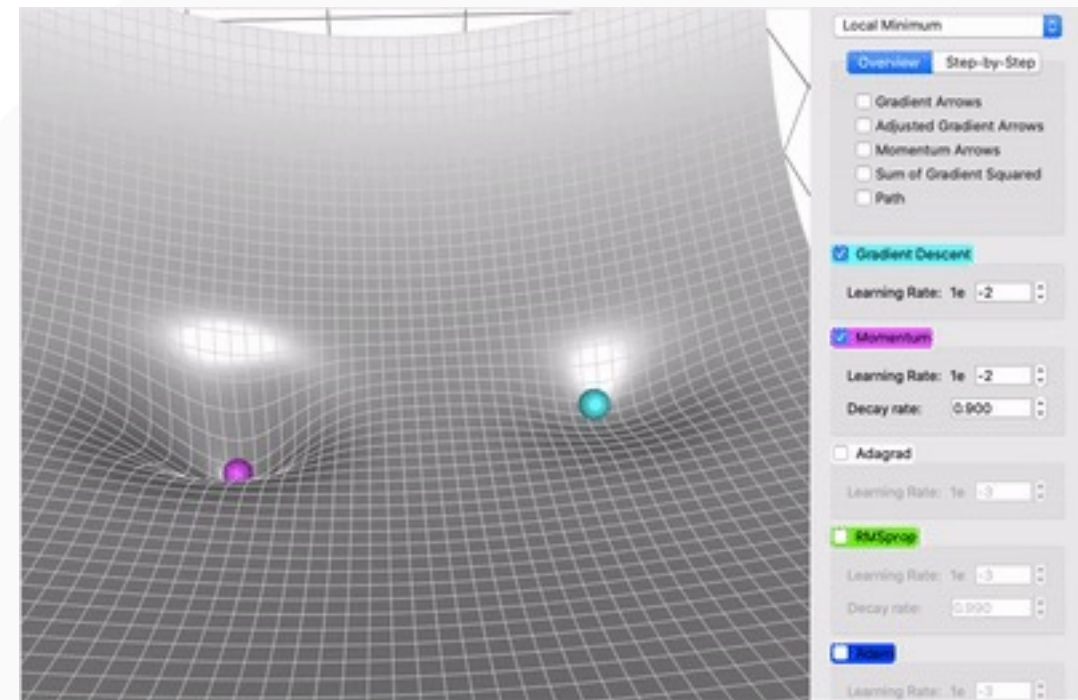


Conventional SGD

$$\mathbf{w} = (w_1, \dots, w_n)$$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla L(\mathbf{w})$$

- Momentum (purple ball) helps not to fall into local minima



SGD with momentum

$$\mathbf{v}_{t+1} = \rho \mathbf{v}_t + \nabla L(\mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \mathbf{v}_{t+1}$$

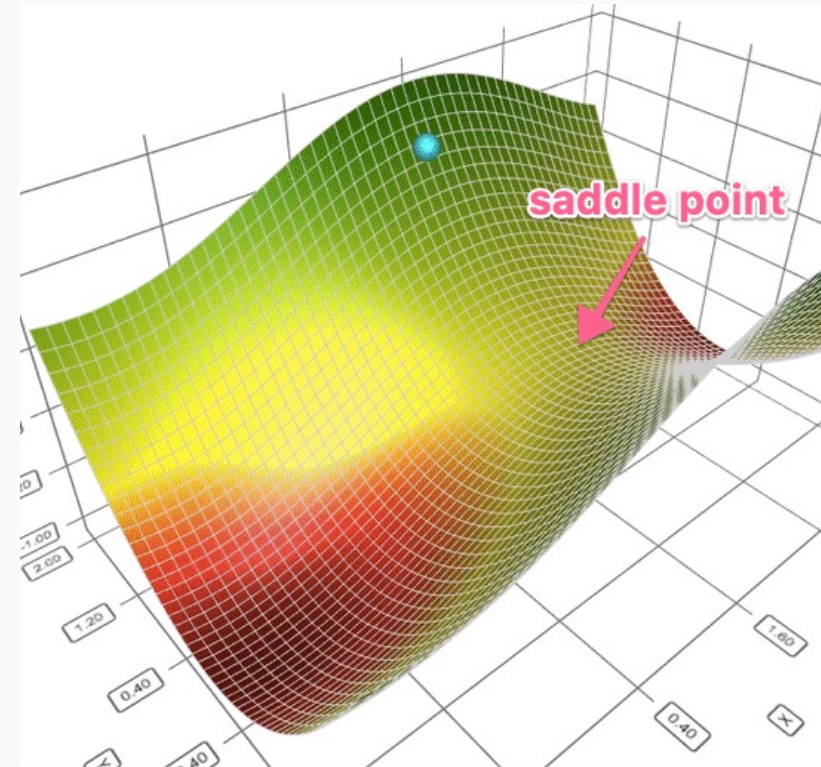
Adaptive Gradient (AdaGrad)

AdaGrad

$$g_{t+1,i} = g_{t,i} + \nabla L(w_i)^2$$

$$w_i \leftarrow w_i - \eta \frac{\nabla L(w_i)}{\sqrt{g_{t+1,i} + 1e^{-5}}}$$

- AdaGrad (Gray ball):
 - Adaptively change the learning rate
 - Prevent excessive moving only in one direction (cyan ball) by increasing the denominator by accumulating all the movement so far)
 - Accelerate the learning from the sparsely acquired movement.
 - Learning gets slower as time goes by -> RMSProp proposed



Momentum (v)

SGD with momentum

$$v_{t+1} = \rho v_t + \nabla L(w)$$

$$w \leftarrow w - \eta v_{t+1}$$

RMSProp

$$g_{t+1} = \rho g_t + (1 - \rho) \nabla L(w)^2$$

$$w \leftarrow w - \eta \frac{\nabla L(w)}{\sqrt{g_{t+1}} + 1e^{-5}}$$

Adam

$$m_{t+1} = \rho_1 m_t + (1 - \rho_1) \nabla L(w) \quad \text{momentum}$$

$$v_{t+1} = \rho_2 v_t + (1 - \rho_2) \nabla L(w)^2 \quad \text{RMSProp}$$

$$w \leftarrow w - \eta \frac{\nabla L(w)}{\sqrt{v_{t+1}} + 1e^{-5}} m_{t+1}$$

Perceptron Training with Gradient Descent (Example2)